

Javier Palomares Homework 1 b

1. Correlations

- When given a data matrix, an easy way to tell if any two columns are correlated is to look at a scatter plot of each column against each other column. For a warm up, do this: Look at the data in DF1 in HW1b Data.zip.

```
In [21]: df1FilePath="./DF1"
df2FilePath="./DF2"

import pandas as pd
df1 = pd.read_csv(df1FilePath,index_col=0)
df2 = pd.read_csv(df2FilePath,index_col=0)
df1.info()
df1.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 0 to 9999
Data columns (total 4 columns):
0      10000 non-null float64
1      10000 non-null float64
2      10000 non-null float64
3      10000 non-null float64
dtypes: float64(4)
memory usage: 390.6 KB
```

Out[21]:

| | 0 | 1 | 2 | 3 |
|--------------|--------------|--------------|--------------|--------------|
| count | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 |
| mean | 0.007704 | 0.009313 | 0.007586 | -0.008455 |
| std | 1.000779 | 1.002686 | 1.000794 | 1.002581 |
| min | -3.471566 | -4.056024 | -3.524182 | -3.930215 |
| 25% | -0.663449 | -0.663886 | -0.666489 | -0.697856 |
| 50% | 0.017736 | 0.009027 | 0.012372 | -0.007986 |
| 75% | 0.667565 | 0.695943 | 0.672933 | 0.668582 |
| max | 3.854101 | 3.908736 | 3.608846 | 3.985592 |

Which columns are (pairwise) correlated? Figure out how to do this with Pandas, and also how to do this with Seaborn.

In [37]: `df1.corr()`

Out[37]:

| | 0 | 1 | 2 | 3 |
|---|-----------|-----------|-----------|-----------|
| 0 | 1.000000 | -0.003998 | 0.990066 | 0.004111 |
| 1 | -0.003998 | 1.000000 | -0.004085 | -0.990235 |
| 2 | 0.990066 | -0.004085 | 1.000000 | 0.004067 |
| 3 | 0.004111 | -0.990235 | 0.004067 | 1.000000 |


Using pandas, we see that columns at position 0 and position 2 are correlated, and columns at position 1 and 3 are negatively correlated

```
In [38]: import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
sns.set(style='white')
corr1 = df1.corr()

# Generate a mask for the upper triangle
mask = np.zeros_like(corr1, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
# Set up the matplotlib figure
f, ax = plt.subplots();

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)
# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr1, mask=mask, cmap=cmap, vmax=1.0, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
```

Out[38]: `<matplotlib.axes._subplots.AxesSubplot at 0x7f64166a2a58>`



Using seaborn, I can see the correlation between columns 0 and 2, and the negative correlation between columns 1 and 3

- Compute the covariance matrix of the data. Write the explicit expression for what this is, and then use any command you like (e.g., `np.cov`) to compute the 4×4 matrix. Explain why the numbers that you get fit with the plots you got.

Let DF1 have size $n \times n$. Let x_i be the i -th column of DF1 and $x_i[j]$ be the j -th component of x_i . Note that columns of DF1 have n components. Let μ_i be the mean of x_i and let σ^2 be the variance of x_i . Then the i, j entry of the covariance matrix is given by:

$$\text{cov}_{ij} = \frac{1}{n} \sum_{k=1}^n (x_i[k] - \mu_i)(x_j[k] - \mu_j)$$

In [39]: `df1.cov()`

Out[39]:

| | 0 | 1 | 2 | 3 |
|---|-----------|-----------|-----------|-----------|
| 0 | 1.001558 | -0.004012 | 0.991624 | 0.004125 |
| 1 | -0.004012 | 1.005378 | -0.004099 | -0.995457 |
| 2 | 0.991624 | -0.004099 | 1.001589 | 0.004081 |
| 3 | 0.004125 | -0.995457 | 0.004081 | 1.005168 |

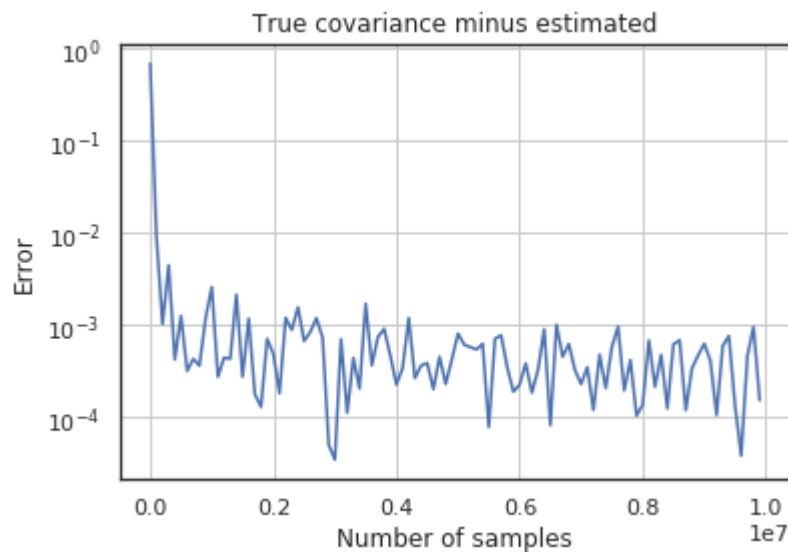
The values in the covariance matrix are consistent with the values in the correlation as they differ by a normalization factor.

- The above problem in reverse. Generate a zero-mean multivariate Gaussian random variable in 3 dimensions, $Z = (X1, X2, X3)$ so that $(X1, X2)$ and $(X1, X3)$ are uncorrelated, but $(X2, X3)$ are correlated. Specifically: choose a covariance matrix that has the above correlations structure, and write this down. Then find a way to generate samples from this Gaussian. Choose one of the non-zero covariance terms (C_{ij} , if C denotes your covariance matrix) and plot it vs the estimated covariance term, as the number of samples you use scales. The goal is to get a visual representation of how the empirical covariance converges to the true (or family) covariance.

```
In [79]: def estimateGaussianCov(mean,cov,numSamples,i,j):
    x = np.random.multivariate_normal(mean, cov,numSamples).T
    estCov = np.cov(x)
    c_ij = estCov[i][j]
    error = c_ij - cov[i][j]
    return error

mean = [0,0,0]
cov = [[1,0,0],[0,1,.99],[0,.99,1]]
#c_1,2 is non zero
i = 1
j = 2
numSamples = list(range(10,10000000,100000))
err = np.zeros(len(numSamples))
for k in range(len(numSamples)):
    n = numSamples[k]
    err[k] = estimateGaussianCov(mean,cov,n,i,j)
```

```
In [85]: err = abs(err)
plt.semilogy(numSamples,err)
plt.title("True covariance minus estimated")
plt.grid(True)
plt.xlabel("Number of samples")
plt.ylabel("Absolute value of error")
plt.show()
```



2. Outliers

Consider the two-dimensional data in DF2 in HW1b Data.zip. Look at a scatter plot of the data. It contains two points that look like potential outliers. Which one is “more” outlying? Propose a transformation of the data that makes it clear that the point at $(-1, 1)$ is more outlying than the point at $(5.5, 5)$, even though the latter point is “farther away” from the nearest points. Plot the data again after performing this transformation. Provide discussion as appropriate to justify your choice of transformation. Hint: if \mathbf{y} comes from a standard Gaussian in two dimensions (i.e., with covariance equal to the two by two identity matrix), and

$$Q = \begin{pmatrix} 2 & \frac{1}{2} \\ \frac{1}{2} & 2 \end{pmatrix}$$

what is the covariance matrix of the random variable $\mathbf{z} = Q\mathbf{y}$? If you are given \mathbf{z} , how would you create a random Gaussian vector with covariance equal to the identity, using \mathbf{z} ?

```
In [82]: df2FilePath="./DF2"

import pandas as pd
df2 = pd.read_csv(df2FilePath,index_col=0)
df2.info()
df2.describe()
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 0 to 9999
Data columns (total 2 columns):
0      10000 non-null float64
1      10000 non-null float64
dtypes: float64(2)
memory usage: 234.4 KB
```

Out[82]:

| | 0 | 1 |
|-------|--------------|--------------|
| count | 10000.000000 | 10000.000000 |
| mean | 0.008139 | 0.008169 |
| std | 1.002321 | 1.002078 |
| min | -3.471566 | -3.524182 |
| 25% | -0.663686 | -0.666489 |
| 50% | 0.017736 | 0.012527 |
| 75% | 0.667909 | 0.673805 |
| max | 5.500000 | 5.000000 |

```
In [83]: df2.plot.scatter(x=0,y=1,c='DarkBlue')
```

```
Out[83]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd108216da0>
```



if $\mathbf{z} = Q\mathbf{y}$, then

$$\text{Covar}(\mathbf{z}) = Q^T \text{Covar}(\mathbf{y}) Q$$

and since $\text{Covar}(\mathbf{y}) = \mathbf{I}$ (identity matrix) then $\text{Covar}(\mathbf{z}) = Q^T Q$.

So by finding the covariance of \mathbf{z} , we have $Q^T Q$.

Q can then be found using Cholesky decomposition. Finally, we can get \mathbf{y} by

$$\mathbf{y} = Q^{-1} \mathbf{z}$$


```
In [93]: covariance = df2.cov()
Q = np.linalg.cholesky(covariance)
Q_inv = np.linalg.inv(Q)
print(Q_inv)
```

```
[[ 0.99768418  0.        ]
 [-6.96054618  7.03339266]]
```

```
In [113]: outlier1 = np.array([-1,1])
outlier2 = np.array([5.5,5])
print(np.dot(Q_inv,outlier1.T).T)
print(np.dot(Q_inv,outlier2.T).T)
```

```
[-0.99768418 13.99393883]
[ 5.48726302 -3.1160407 ]
```

```
In [114]: # the transformed data
z = np.dot(Q_inv, df2.T)
z = z.T
plt.scatter(z[:,0], z[:,1])
plt.show()
```



We can see that the point at (-1,1) has been mapped to a point further from the origin than (5.5,5) was mapped to.

3. Even More Standard Error

(This is to be completed only after you've completed the last written exercise below). In one of the written exercises below, you derive an expression for what is called the Standard Error: where β denotes the “truth,” $\hat{\beta}$ denotes the value we compute using least squares linear regression, and Z and e are as in the exercise below, you find: $\hat{\beta} - \beta = Ze$. If we know the distribution of the noise (the distribution generating the noise vectors, e_i), then we know the distribution for the error, $(\hat{\beta} - \beta)$. This allows us to answer the question given in class: if we solve a regression and obtain value $\hat{\beta}$, how can we tell if it is statistically significant? The answer is: we compare the size of $\hat{\beta}$ to the spread introduced by the noise (i.e., the standard error), and we ask: what is the likelihood that the true $\beta = 0$, and what we observed was purely due to the noise. If the noise is Gaussian (normal), i.e., $e_i \sim N(0, \sigma^2)$, and if the values of the x_i are normalized, then we expect error of the size $\frac{\sigma}{\sqrt{n}}$, as this is roughly the standard deviation of the expression for the error that you derive above. This means: if you have twice the data points, you should expect the error to be reduced by about 1.4 (the formula says that the standard deviation of the error would decrease by a factor of $1/\sqrt{2}$).

Compute this empirically, as follows: We will generate data for a regression problem, solve it, and see what the error is: Generate data as I did in the example from class: $x_i \sim N(0, 1)$, $e_i \sim N(0, 1)$. Generate y by $y_i = \beta_0 + x_i\beta + e_i$, where $\beta_0 = -3$ and $\beta = 0$. Note that since $\beta = 0$, this means that y and x are unrelated! The question we are exploring here is as follows: when we solve a regression problem, we are not going to find $\hat{\beta} = 0$ - we will find that $\hat{\beta}$ takes some other values, hopefully close to zero. How do we know if the value of $\hat{\beta}$ we get is statistically meaningful?

```
In [118]: beta0 = -3
          beta = 0
          meanX = 0
          varX = 1
          meanE = 0
          varE = 1
```

- By creating fresh data and each time computing $\hat{\beta}$ and recording $\hat{\beta} - \beta$, compute the empirical standard deviation of the error for $n = 150$ (the number we used in class). In class, in the exercise where I tried to find a linear regression of y vs. noise, we found $\hat{\beta} = -0.15$. Given your empirical computation of the standard deviation of the error, how significant is the value -0.15 ?

```
In [124]: import numpy as np
          import math
          def computeBetaHat(meanX,varX,meanE,varE,n):
              # get x and e from an normal distribution
              x = np.random.normal(meanX,math.sqrt(varX),n)
              e = np.random.normal(meanE,math.sqrt(varE),n)
              # compute y from x and e
              y = beta0 + beta*x + e
              # expression computed from last problem of homework
              betaHat = np.dot(x,y) / np.dot(x,x)
              return betaHat

          n = 150
          betaHat = computeBetaHat(meanX,varX,meanE,varE,n)
          print("The computed value of beta (betaHat) is {} for n={}".format(betaHat,n))
```

The computed value of beta (betaHat) is -0.2548654898775995 for n=150

```
In [125]: def meanStdDevBeta(numTrials,meanX,varX,meanE,varE,n):
          betaSamples = np.zeros(numTrials)
          for i in range(numTrials):
              betaSamples[i] = computeBetaHat(meanX,varX,meanE,varE,n)
          meanBeta = np.mean(betaSamples)
          stdBeta = np.std(betaSamples)
          return meanBeta,stdBeta
```

```
In [126]: # compute the standard deviation for the computed value of beta^ when
          n = 150
          numTrials = 1000;
          meanBeta,stdBeta = meanStdDevBeta(numTrials,meanX,varX,meanE,varE,n)
          print("After {} trials, betaHat has a mean of {} with a standard deviation of {}".format(numTrials,meanBeta,stdBeta))
```

After 1000 trials, betaHat has a mean of -0.0009606705594536155 with a standard deviation of 0.2516646103140726

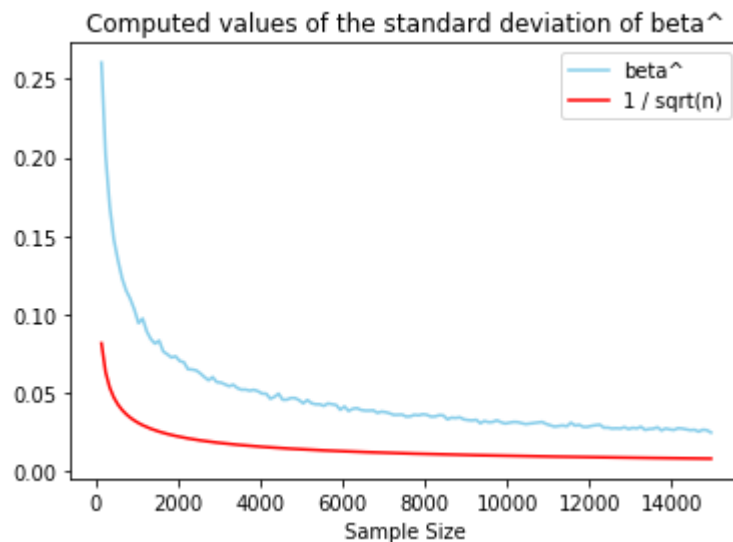
Finding a value for $\hat{\beta} = -0.15$ is about half a standard deviation away from the mean (using the empirical values)

- Now repeat the above experiment for different values of n . Plot these values, and on the same plot, plot $1/\sqrt{n}$. How is the fit?

```
In [129]: sampleSizes = np.arange(150, 15000,100)
betaHatMean = np.zeros(len(sampleSizes))
betaHatStd = np.zeros(len(sampleSizes))
# compute the mean and stdDev for betaHat for different values of n
for i in range(len(sampleSizes)):
    betaHatMean[i],betaHatStd[i] = meanStdDevBeta(numTrials,meanX,varX,meanE,varE,sampleSizes[i])
```

```
In [130]: import matplotlib.pyplot as plt
plt.plot(sampleSizes,betaHatStd,color='skyblue',label = "beta^")
plt.plot(sampleSizes,1.0/np.sqrt(sampleSizes),color='red',label = '1 / sqrt(n)')
plt.title("Computed values of the standard deviation of beta^")
plt.xlabel("Sample Size")
plt.legend()
```

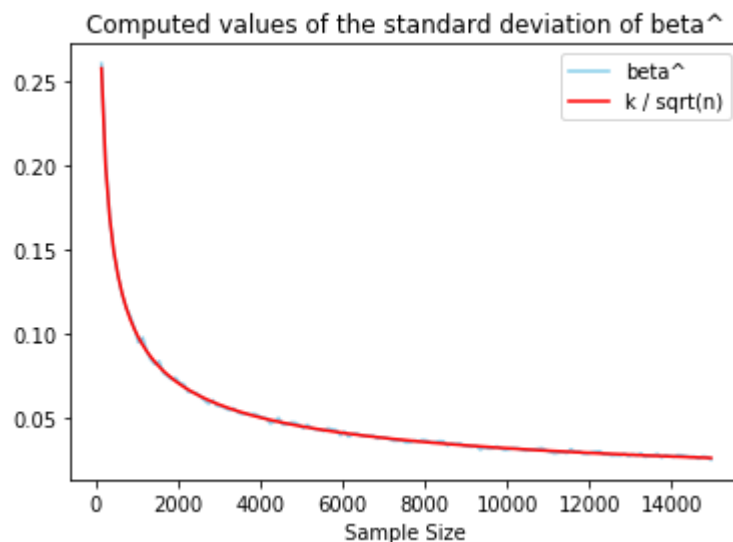
Out[130]: <matplotlib.legend.Legend at 0x7fd106103208>



```
In [134]: sigma = np.divide(betaHatStd,1.0/np.sqrt(sampleSizes))
k = np.mean(sigma)
```

```
In [136]: import matplotlib.pyplot as plt
plt.plot(sampleSizes,betaHatStd,color='skyblue',label = "beta^")
plt.plot(sampleSizes,k/np.sqrt(sampleSizes),color='red',label = 'k /
sqrt(n)')
plt.title("Computed values of the standard deviation of beta^")
plt.xlabel("Sample Size")
plt.legend()
```

Out[136]: <matplotlib.legend.Legend at 0x7fd10eff77f0>



$\frac{1}{\sqrt{n}}$ is off from the empirical value of the standard deviation of $\hat{\beta}$ by a constant multiple. If I instead plot $\frac{k}{\sqrt{n}}$ where k is the mean of $\frac{std(\hat{\beta})}{1.0/\sqrt{n}}$, then the 2 lines are almost on top of each other.

Names and Frequencies

The goal of this exercise is for you to get more experience with Pandas, and to get a chance to explore a cool data set. Download the file Names.zip from Canvas. This contains the frequency of all names that appeared more than 5 times on a social security application from 1880 through 2015

- Write a program that on input k and XXXX, returns the top k names from year XXXX.

```
In [1]: import pandas as pd
def topKNames(k,year):
    # read the file containing the name frequency for the year
    filepath = "./Names/yob{}.txt".format(year)
    df = pd.read_csv(filepath,header=None)
    # first column of the dataframe is the name, followed by the gender and last by the frequency
    df.columns = ['name','gender','frequency']
    # sort the dataframe by the frequency in descending order
    df= df.sort_values(by='frequency',ascending=False)
    # return the name of the first k rows of the data
    topKNames = df.iloc[0:k,0]
    return topKNames.to_list()
```

```
In [62]: print(topKNames(7,1990))
```

```
['Michael', 'Christopher', 'Jessica', 'Ashley', 'Matthew', 'Joshua', 'Brittany']
```

- Write a program that on input Name returns the frequency for men and women of the name *Name*.

```
In [101]: # find the frequency of a name in a year for men and women
def freqMFNameByYear(name,year):
    # read the file containing the name frequency for the year
    filepath = "./Names/yob{}.txt".format(year)
    df = pd.read_csv(filepath,header=None)
    # first column of the dataframe is the name, followed by the gender and last by the frequency
    df.columns = ['name','gender','frequency']
    # find the rows for the give name using case insensitive comparison
    rows = df.loc[df['name'].str.lower() == name.lower()]
    # iterate over the rows
    menFrequency = 0
    womenFrequency = 0
    for index,row in rows.iterrows():
        if(row['gender'].lower() == 'm'):
            menFrequency += row['frequency']
        elif(row['gender'].lower() == 'f'):
            womenFrequency += row['frequency']
    return menFrequency,womenFrequency
```

```
In [105]: men,women = freqMFNameByYear("Ashley",1990)
print("{} men were named Ashley in 1990. {} women were named Ashley in 1990".format(men,women))
```

```
239 men were named Ashley in 1990. 45553 women were named Ashley in 1990
```

- It could be that names are more diverse now than they were in 1880, so that a name may be relatively the most popular, though its frequency may have been decreasing over the years. Modify the above to return the relative frequency. Note that in the next coming lectures we will learn how to quantify diversity using entropy.

```
In [136]: def getTotalMWFreqByYear(year):
            filepath = "./Names/yob{}.txt".format(year)
            df = pd.read_csv(filepath, header=None)
            # first column of the dataframe is the name, followed by the gen
der and last by the frequency
            df.columns = ['name', 'gender', 'frequency']
            # find the total frequency for men and women
            menTotal = df.loc[df['gender'].str.lower() == "m"]['frequency'].sum()
            womenTotal = df.loc[df['gender'].str.lower() == "f"]['frequency'].sum()
            return menTotal, womenTotal

            # find the relative frequency of a name in a year for men and women
            def relFreqMFNameByYear(name, year):
                # get the frequency of a name in a year for men and women
                men, women = freqMFNameByYear(name, year)
                menTotal, womenTotal = getTotalMWFreqByYear(year)
                return men/menTotal, women/womenTotal
```

```
In [137]: men, women = relFreqMFNameByYear("Ashley", 1990)
            print("{}% of men were named Ashley in 1990. {}% of women were named Ashley in 1990".format(men*100, women*100))

0.011643417182176121% of men were named Ashley in 1990. 2.40029676358028% of women were named Ashley in 1990
```

- Find all the names that used to be more popular for one gender, but then became more popular for another gender.

```
In [4]: startYear = 1880
endYear = 2015
years = np.arange(startYear,endYear+1)
# initialize a dictionary for men and women to hold the count of a name per year
namesM = {}
namesW = {}
columns = ['name', 'gender', 'frequency']
for year in years:
    filepath = "./Names/yob{}.txt".format(year)
    df = pd.read_csv(filepath, header=None)
    # first column of the dataframe is the name, followed by the gender and last by the frequency
    df.columns = columns
    for index, row in df.iterrows():
        name = row['name']
        isFemale = row['gender'].lower() == 'f'
        frequency = row['frequency']
        if isFemale:
            names = namesW
        else:
            names = namesM
        counts = names.get(name, {})
        # store the frequency in the counts
        counts[year] = frequency
        names[name] = counts
```

```

In [6]: import numpy as np
# Note: there are thousands of names who flip. For pdf purposes, I'll
# only show 20 flips for men and women.
numLines = 20;
i = 0;
# now iterate over all the names and find if a name becomes more popular
# for one gender vs the other
for name, countsM in namesM.items():
    countsW = namesW.get(name, {})
    if len(countsW) == 0:
        # no women are named name
        continue
    # there are women with the name
    sign = None
    for year in years:
        numMen = countsM.get(year, 0)
        numWomen = countsW.get(year, 0)
        diff = numMen - numWomen
        s = np.sign(diff)
        if not sign:
            sign = s
            continue
        if (s != sign):
            i+=1
            sign = s
            if (s > 0):
                if (i < numLines):
                    print("{} become more popular for men({}) that women({}) in {}".format(name, numMen, numWomen, year))
            elif (s < 0):
                if (i < numLines):
                    print("{} become more popular for women({}) that men({}) in {}".format(name, numWomen, numMen, year))
            else:
                if (i < numLines):
                    print("{} became as popular for men({}) and women({}) in {}".format(name, numMen, numWomen, year))

i=0
# do the same for women's name
for name, countsW in namesW.items():
    countsM = namesM.get(name, {})
    if len(countsM) == 0:
        # no men are named name
        continue
    # there are men with the name
    sign = None
    for year in years:
        numMen = countsM.get(year, 0)
        numWomen = countsW.get(year, 0)
        diff = numMen - numWomen
        s = np.sign(diff)
        if not sign:
            sign = s
            continue
        if (s != sign):

```

```

sign = s
i=i+1
if (s > 0):
    if( i < numLines):
        print("{} become more popular for men({}) that wo
men({}) in {}".format(name,numMen,numWomen,year))
    elif(s < 0):
        if( i< numLines):
            print("{} become more popular for women({}) that
men({}) in {}".format(name,numWomen,numMen,year))
        else:
            if(i < numLines):
                print("{} became as popular for men({}) and women
({}) in {}".format(name,numMen,numWomen,year))

```

Ed became as popular for men(0) and women(0) in 2012
Charley become more popular for women(63) that men(61) in 1987
Charley became as popular for men(56) and women(56) in 1990
Marion become more popular for women(184) that men(163) in 1883
Marion become more popular for men(184) that women(181) in 1884
Marion become more popular for women(198) that men(141) in 1885
Marion become more popular for men(299) that women(293) in 1972
Marion became as popular for men(229) and women(229) in 1977
Marion become more popular for women(211) that men(187) in 1981
Marion become more popular for men(169) that women(165) in 1984
Marion become more popular for women(168) that men(154) in 1985
Marion become more popular for men(164) that women(162) in 1986
Marion become more popular for women(150) that men(137) in 1987
Marion become more popular for men(117) that women(115) in 1996
Marion become more popular for women(141) that men(102) in 2000
Marion become more popular for men(158) that women(119) in 2003
Marion become more popular for women(111) that men(102) in 2011
Marion became as popular for men(109) and women(109) in 2012
Jessie become more popular for men(1094) that women(1070) in 1948
Jessie become more popular for men(1094) that women(1070) in 1948
Jessie become more popular for women(1031) that men(1023) in 1949
Jessie become more popular for men(1019) that women(937) in 1950
Jessie become more popular for women(1276) that men(1117) in 1986
Jessie become more popular for men(1255) that women(1116) in 1990
Jessie become more popular for women(1096) that men(1048) in 1992
Jessie become more popular for men(627) that women(618) in 2003
Jessie become more popular for women(607) that men(570) in 2004
Jessie become more popular for men(590) that women(507) in 2005
Jessie become more popular for women(493) that men(488) in 2008
Jessie become more popular for men(459) that women(421) in 2009
Jessie become more popular for women(419) that men(406) in 2010
Myrtle became as popular for men(0) and women(0) in 1998
Myrtle became as popular for men(0) and women(0) in 2006
Myrtle became as popular for men(0) and women(0) in 2015
Blanche became as popular for men(0) and women(0) in 2015
Maud became as popular for men(0) and women(0) in 1963
Maud became as popular for men(0) and women(0) in 1971
Maud became as popular for men(0) and women(0) in 1975

Visualization Tools and Missing/Hidden Values

Visualization is important both for exploring the data, as well as for explaining what you have done. There are a huge number of such tools now available. This exercise walks through various functionalities of matplotlib and pandas.

- The first part of this exercise was created by Dataquest. Run through the commands given in this tutorial: <https://www.dataquest.io/blog/matplotlib-tutorial/> (<https://www.dataquest.io/blog/matplotlib-tutorial/>) and understand the code.

Exploring tweets with Pandas

- Import the Pandas library.
- Read tweets.csv into a Pandas DataFrame.
- Print the first 5 rows of the DataFrame

```
In [44]: import pandas as pd
tweets = pd.read_csv("tweets.csv")
tweets.head()
```

Out[44]:

| | id | id_str | user_location | user_bg_color | retweet_count | user_name | polarity |
|---|----|--------------------|--------------------|---------------|---------------|--------------|----------|
| 0 | 1 | 729828033092149248 | Wheeling WV | 022330 | 0 | Jaybo26003 | 0.00 |
| 1 | 2 | 729828033092161537 | NaN | C0DEED | 0 | britttany_ns | 0.15 |
| 2 | 3 | 729828033566224384 | NaN | C0DEED | 0 | JeffriesLori | 0.00 |
| 3 | 4 | 729828033893302272 | global | C0DEED | 0 | WhorunsGOVs | 0.00 |
| 4 | 5 | 729828034178482177 | California, USA | 131516 | 0 | BJCG0830 | 0.00 |

Generating a candidates column

- Create a function that finds what candidate names occur in a piece of text.
- Use the apply method on DataFrames to generate a new column called candidate that contains what candidate(s) the tweet mentions.

```
In [140]: def get_candidate(row):  
           candidates = []  
           text = row["text"].lower()  
           if "clinton" in text or "hillary" in text:  
               candidates.append("clinton")  
           if "trump" in text or "donald" in text:  
               candidates.append("trump")  
           if "sanders" in text or "bernie" in text:  
               candidates.append("sanders")  
           return ", ".join(candidates)  
  
           tweets["candidate"] = tweets.apply(get_candidate,axis=1)
```

Making the first plot

Importing matplotlib


```
In [141]: import matplotlib.pyplot as plt  
           import numpy as np  
           %matplotlib inline
```

Making a bar plot

- Use the value_counts method on Pandas Series to count up how many tweets mention each candidate.
- Use plt.bar to create a bar plot. We'll pass in a list of numbers from 0 to the number of unique values in the candidate column as the x-axis input, and the counts as the y-axis input.
- Display the counts so we have more context about what each bar represents.

```
In [142]: counts = tweets["candidate"].value_counts()
plt.bar(range(len(counts)), counts)
plt.show()

print(counts)
```



```
trump          119998
clinton, trump 30521
               25429
sanderson      25351
clinton        22746
clinton, sanders 6044
clinton, trump, sanders 4219
trump, sanderson 3172
Name: candidate, dtype: int64
```


Customizing plots

- Convert the created and user_created columns to the Pandas datetime type.
- Create a user_age column that is the number of days since the account was created.
- Create a histogram of user ages.
- Show the histogram.

```
In [143]: from datetime import datetime

tweets["created"] = pd.to_datetime(tweets["created"])
tweets["user_created"] = pd.to_datetime(tweets["user_created"])


tweets["user_age"] = tweets["user_created"].apply(lambda x: (datetime
.now() - x).total_seconds() / 3600 / 24 / 365)
plt.hist(tweets["user_age"])
plt.show()
```



Adding labels

- Generate the same histogram we did before.
- Draw a title onto the histogram.
- Draw an x axis label onto the histogram.
- Draw a y axis label onto the histogram.
- Show the plot.


```
In [144]: plt.hist(tweets["user_age"])
plt.title("Tweets mentioning candidates")
plt.xlabel("Twitter account age in years")
plt.ylabel("# of tweets")
plt.show()
```



Making a stacked histogram

- Generate three Pandas series, each containing the user_age data only for tweets about a certain candidate.
- Make a stacked histogram by calling the hist method with additional options.
 - Specifying a list as the input will plot three sets of histogram bars.
 - Specifying stacked=True will stack the three sets of bars.
 - Adding the label option will generate the correct labels for the legend.
- Call the plt.legend method to draw a legend in the top right corner.
- Add a title, x axis, and y axis labels.
- Show the plot.

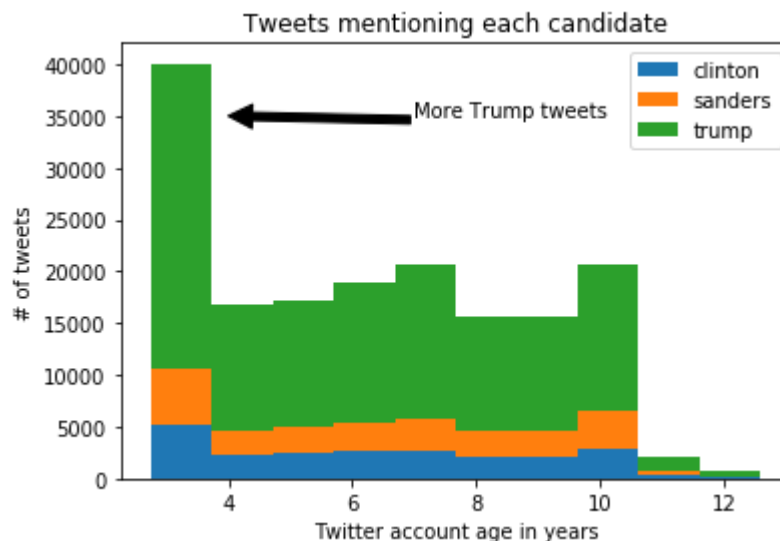
```
In [146]: cl_tweets = tweets["user_age"][tweets["candidate"] == "clinton"]
sa_tweets = tweets["user_age"][tweets["candidate"] == "sanders"]
tr_tweets = tweets["user_age"][tweets["candidate"] == "trump"]
plt.hist([
    cl_tweets,
    sa_tweets,
    tr_tweets
],
    stacked=True,
    label=["clinton", "sanders", "trump"]
)
plt.legend()
plt.title("Tweets mentioning each candidate")
plt.xlabel("Twitter account age in years")
plt.ylabel("# of tweets")
plt.show()
```



Annotating the histogram

In the code below, we'll make the same histogram as we did above, but we'll call the `plt.annotate` method to add an annotation to the plot.

```
In [167]: plt.hist([
            cl_tweets,
            sa_tweets,
            tr_tweets
        ],
        stacked=True,
        label=["clinton", "sanders", "trump"]
    )
plt.legend()
plt.title("Tweets mentioning each candidate")
plt.xlabel("Twitter account age in years")
plt.ylabel("# of tweets")
plt.annotate('More Trump tweets', xy=(4, 35000), xytext=(7, 35000),
            arrowprops=dict(facecolor='black'))
plt.show()
```



Multiple subplots

Extracting colors

- Use the apply method to go through each row in the user_bg_color column, and extract how much red is in it.
- Use the apply method to go through each row in the user_bg_color column, and extract how much blue is in it.

```
In [168]: import matplotlib.colors as colors

tweets["red"] = tweets["user_bg_color"].apply(lambda x: colors.hex2color('#{0}'.format(x))[0])
tweets["blue"] = tweets["user_bg_color"].apply(lambda x: colors.hex2color('#{0}'.format(x))[2])
```

Creating the plot

- Generate a Figure and multiple Axes with the subplots method. The axes will be returned as an array.
- The axes are returned in a 2x2 NumPy array. We extract each individual Axes object by using the flat property of arrays. This gives us 4 Axes objects we can work with.
- Plot a histogram in the first Axes using the hist method.
- Set the title of the first Axes to Red in all backgrounds using the set_title method. This performs the same function as plt.title.
- Plot a histogram in the second Axes using the hist method.
- Set the title of the second Axes to Red in Trump tweeters using the set_title method.
- Plot a histogram in the third Axes using the hist method.
- Set the title of the third Axes to Blue in all backgrounds using the set_title method. This performs the same function as plt.title.
- Plot a histogram in the fourth Axes using the hist method.
- Set the title of the fourth Axes to Blue in Trump tweeters using the set_title method.
- Call the plt.tight_layout method to reduce padding in the graphs and fit all the elements.
- Show the plot.

```
In [169]: fig, axes = plt.subplots(nrows=2, ncols=2)
ax0, ax1, ax2, ax3 = axes.flat


ax0.hist(tweets["red"])
ax0.set_title('Red in backgrounds')

ax1.hist(tweets["red"][tweets["candidate"] == "trump"].values)
ax1.set_title('Red in Trump tweeters')

ax2.hist(tweets["blue"])
ax2.set_title('Blue in backgrounds')

ax3.hist(tweets["blue"][tweets["candidate"] == "trump"].values)
ax3.set_title('Blue in Trump tweeters')

plt.tight_layout()
plt.show()
```



Removing common background colors

Here's how to find the most common colors in background colors:


```
In [171]: tweets["user_bg_color"].value_counts()
```

```

Out[171]: C0DEED      108977
          000000      31119
          F5F8FA      25597
          131516      7731
          1A1B1F      5059
          022330      4300
          0099B9      3958
          642D8B      3767
          FFFFFFFF     3101
          9AE4E8      2651
          ACDED6      2383
          352726      2338
          C6E2EE      1978
          709397      1518
          EBEbeb      1475
          FF6699      1370
          BADFCD      1336
          FFF04D      1300
          EDECE9      1225
          B2DFDA      1218
          DBE9ED      1113
          ABB8C2      1101
          8B542B      1073
          3B94D9       623
          89C9FA       414
          DD2E44       351
          94D487       318
          4A913C       300
          9266CC       287
          F5ABB5       267
          ...
          BAE0F5        1
          FBFBFB        1
          F0488B        1
          D0DBE0        1
          878287        1
          13BEED        1
          233B47        1
          8A2D5F        1
          AA86B1        1
          C2CC31        1
          2028BD        1
          332524        1
          E50E06        1
          8A868A        1
          EAEBED        1
          FCA9FC        1
          993893        1
          C20E0E        1
          BF6F9C        1
          FF7105        1
          2288CC        1
          F38630        1
          09090A        1
          1F1A1E        1
          D9007E        1
          5590AD        1

```

```

300030      1
0B4C5F      1
CCD6D6      1
8F0E8F      1
Name: user_bg_color, Length: 6970, dtype: int64

```

- Remove C0DEED, 000000, and F5F8FA from user_bg_color.
- Create a function with out plotting logic from the last chart inside.
- Plot the same 4 plots from before without the most common colors in user_bg_color.

```

In [172]: tc = tweets[~tweets["user_bg_color"].isin(["C0DEED", "000000", "F5F8FA"])]

def create_plot(data):
    fig, axes = plt.subplots(nrows=2, ncols=2)
    ax0, ax1, ax2, ax3 = axes.flat

    ax0.hist(data["red"])
    ax0.set_title('Red in backgrounds')

    ax1.hist(data["red"][data["candidate"] == "trump"].values)
    ax1.set_title('Red in Trump tweets')


    ax2.hist(data["blue"])
    ax2.set_title('Blue in backgrounds')

    ax3.hist(data["blue"][data["candidate"] == "trump"].values)
    ax3.set_title('Blue in Trump tweeters')

    plt.tight_layout()
    plt.show()

create_plot(tc)

```



Plotting sentiment

- Group tweets by candidate, and compute the mean and standard deviation for each numerical column (including polarity).
- Create a Figure that's 7 inches by 7 inches, with 2 Axes objects, arranged vertically.
- Create a bar plot of the standard deviation the first Axes object.
 - Set the tick labels using the `set_xticklabels` method, and rotate the labels 45 degrees using the `rotation` argument.
 - Set the title.
- Create a bar plot of the mean on the second Axes object.
 - Set the tick labels.
 - Set the title.
- Show the plot.

```

In [173]: gr = tweets.groupby("candidate").agg([np.mean, np.std])

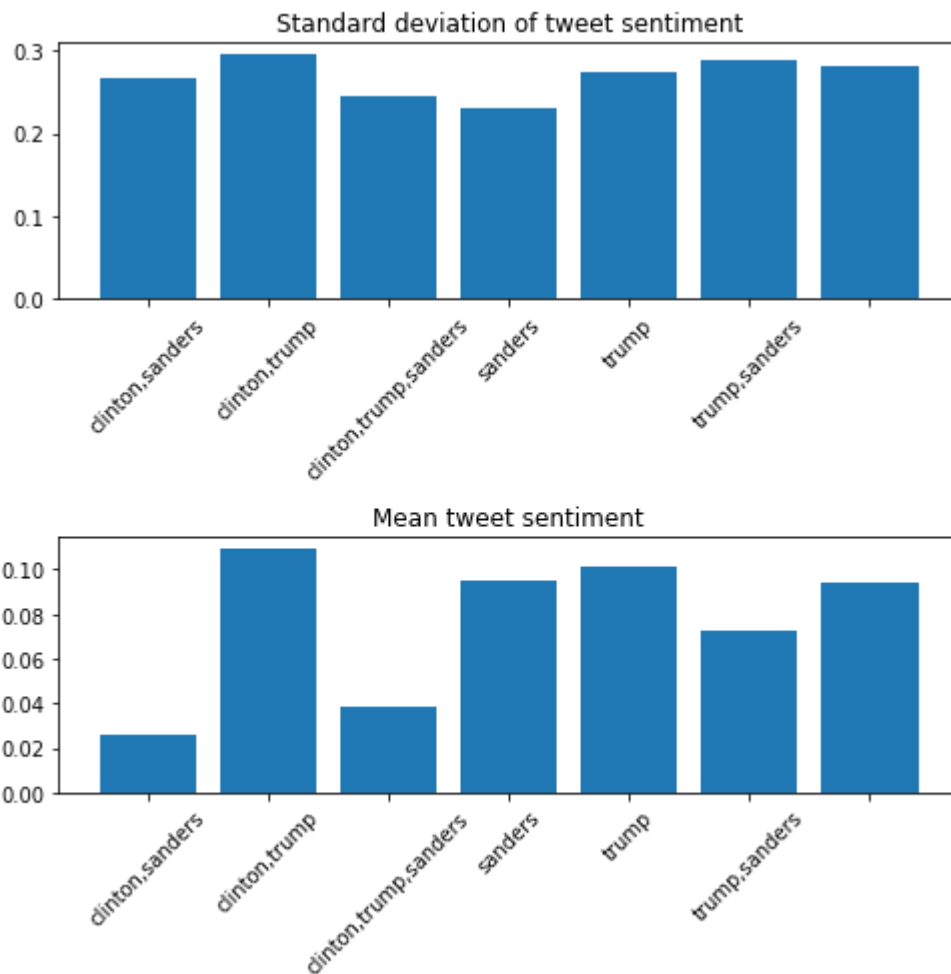
fig, axes = plt.subplots(nrows=2, ncols=1, figsize=(7, 7))
ax0, ax1 = axes.flat

std = gr["polarity"]["std"].iloc[1:]
mean = gr["polarity"]["mean"].iloc[1:]
ax0.bar(range(len(std)), std)
ax0.set_xticklabels(std.index, rotation=45)
ax0.set_title('Standard deviation of tweet sentiment')

ax1.bar(range(len(mean)), mean)
ax1.set_xticklabels(mean.index, rotation=45)
ax1.set_title('Mean tweet sentiment')

plt.tight_layout()
plt.show()

```



Generating a side by side bar plot

Generating tweet lengths

- Define a function to mark a tweet as short if it's less than 100 characters, medium if it's 100 to 135 characters, and long if it's over 135 characters.
- Use apply to generate a new column tweet_length.
- Figure out how many tweets by each candidate fall into each group.


```
In [175]: def tweet_lengths(text):  
            if len(text) < 100:  
                return "short"  
            elif 100 <= len(text) <= 135:  
                return "medium"  
            else:  
                return "long"  
  
            tweets["tweet_length"] = tweets["text"].apply(tweet_lengths)  
  
            tl = {}  
            for candidate in ["clinton", "sanderson", "trump"]:  
                tl[candidate] = tweets["tweet_length"][tweets["candidate"] == candidate].value_counts()
```

Plotting

- Create a Figure and a single Axes object.
- Define the width for each bar, .5.
- Generate a sequence of values, x, that is 0, 2, 4. Each value is the start of a category, such as short, medium, and long. We put a distance of 2 between each category so we have space for multiple bars.
- Plot clinton tweets on the Axes object, with the bars at the positions defined by x.
- Plot sanderson tweets on the Axes object, but add width to x to move the bars to the right.
- Plot trump tweets on the Axes object, but add width * 2 to x to move the bars to the far right.
- Set the axis labels and title.
- Use set_xticks to move the tick labels to the center of each category area.
- Set tick labels.

```
In [177]: fig, ax = plt.subplots()
width = .5
x = np.array(range(0, 6, 2))
ax.bar(x, tl["clinton"], width, color='g')
ax.bar(x + width, tl["sanderson"], width, color='b')
ax.bar(x + (width * 2), tl["trump"], width, color='r')

ax.set_ylabel('# of tweets')
ax.set_title('Number of Tweets per candidate by length')
ax.set_xticks(x + (width * 1.5))
ax.set_xticklabels(('long', 'medium', 'short'))
ax.set_xlabel('Tweet length')
plt.show()
```



- Suppose that you would now like to plot some of the results by state. As you will see, the state information is sometimes missing, and other times it comes in varying forms. Figure out how to aggregate the results by state. The challenge here: how many of the tweets can you (correctly) assign to a state? Note: depending on how well you want to do (i.e., how many tweets you want to correctly assign to their state), this is not an easy problem!

```
In [64]: states = {
    'AK': 'Alaska',
    'AL': 'Alabama',
    'AR': 'Arkansas',
    'AS': 'American Samoa',
    'AZ': 'Arizona',
    'CA': 'California',
    'CO': 'Colorado',
    'CT': 'Connecticut',
    'DC': 'District of Columbia',
    'DE': 'Delaware',
    'FL': 'Florida',
    'GA': 'Georgia',
    'GU': 'Guam',
    'HI': 'Hawaii',
    'IA': 'Iowa',
    'ID': 'Idaho',
    'IL': 'Illinois',
    'IN': 'Indiana',
    'KS': 'Kansas',
    'KY': 'Kentucky',
    'LA': 'Louisiana',
    'MA': 'Massachusetts',
    'MD': 'Maryland',
    'ME': 'Maine',
    'MI': 'Michigan',
    'MN': 'Minnesota',
    'MO': 'Missouri',
    'MP': 'Northern Mariana Islands',
    'MS': 'Mississippi',
    'MT': 'Montana',
    'NA': 'National',
    'NC': 'North Carolina',
    'ND': 'North Dakota',
    'NE': 'Nebraska',
    'NH': 'New Hampshire',
    'NJ': 'New Jersey',
    'NM': 'New Mexico',
    'NV': 'Nevada',
    'NY': 'New York',
    'OH': 'Ohio',
    'OK': 'Oklahoma',
    'OR': 'Oregon',
    'PA': 'Pennsylvania',
    'PR': 'Puerto Rico',
    'RI': 'Rhode Island',
    'SC': 'South Carolina',
    'SD': 'South Dakota',
    'TN': 'Tennessee',
    'TX': 'Texas',
    'UT': 'Utah',
    'VA': 'Virginia',
    'VI': 'Virgin Islands',
    'VT': 'Vermont',
    'WA': 'Washington',
    'WI': 'Wisconsin',
```



```
        'WV': 'West Virginia',
        'WY': 'Wyoming'
    }
import re
def getStateFromLocation(location):
    # search for the state abbreviation or full name in the string
    for key,value in states.items():
        if re.search(r'\b' + key + r'\b',location):
            return value
        elif re.search(r'\b' + value + r'\b',location):
            return value
    # found no state
    return np.nan
```

```
In [42]: import json
import reverse_geocode as rg
import geocoder
import numpy as np

# the coordinates are given as a json string
# parse the string into a dictionary
def getCoordinates(geo):
    g = json.loads(geo)
    return g["coordinates"]

def getStateFromCoord(coord):
    if not coord:
        return np.nan
    location = rg.search(coord)
    state = np.nan
    if location:
        state = location[0]['admin1']
    return state

def getStateFromLoc(loc):
    location = geocoder.arcgis(loc)
    state = np.nan
    if location:
        coord = [location.json['lat'], location.json['lng']]
        state = getStateFromCoord(coord)
    return state

def getState(row):
    state = np.nan
    # try to get the state from a geolocation
    # or from the user's location
    geo = row["geo"]
    loc = row["user_location"]
    # map the geolocation to a state
    if(pd.isna(geo) == False):
        try:
            coord = getCoordinates(geo)
            state = getStateFromCoord(coord)
        except:
            return np.nan
    # map the location to a state
    elif(pd.isna(loc) == False):
        try:
            state = getStateFromLocation(loc)
        except:
            return np.nan
    return state
```

```
In [65]: tweets["state"] = tweets.apply(getState,axis=1)
```

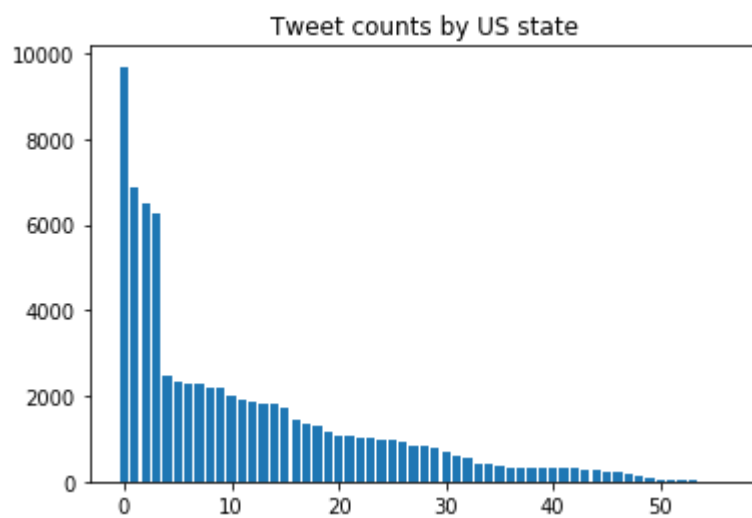
In [66]: tweets.head()

Out[66]:

| | id | id_str | user_location | user_bg_color | retweet_count | user_name | polarity |
|---|----|--------------------|-----------------|---------------|---------------|--------------|----------|
| 0 | 1 | 729828033092149248 | Wheeling WV | 022330 | 0 | Jaybo26003 | 0.00 |
| 1 | 2 | 729828033092161537 | NaN | C0DEED | 0 | britttany_ns | 0.15 |
| 2 | 3 | 729828033566224384 | NaN | C0DEED | 0 | JeffriesLori | 0.00 |
| 3 | 4 | 729828033893302272 | global | C0DEED | 0 | WhorunsGOVs | 0.00 |
| 4 | 5 | 729828034178482177 | California, USA | 131516 | 0 | BJCG0830 | 0.00 |

```
In [71]: counts = tweets["state"].value_counts()
plt.bar(range(len(counts)), counts)
plt.title('Tweet counts by US state')
plt.show()

print(counts)
```



| | |
|--------------------------|------|
| California | 9690 |
| Texas | 6868 |
| New York | 6511 |
| Florida | 6257 |
| Washington | 2472 |
| North Carolina | 2344 |
| Pennsylvania | 2303 |
| District of Columbia | 2278 |
| New Jersey | 2213 |
| Georgia | 2191 |
| Arizona | 2023 |
| Ohio | 1906 |
| Virginia | 1856 |
| Illinois | 1837 |
| Tennessee | 1828 |
| Michigan | 1715 |
| Colorado | 1461 |
| Massachusetts | 1337 |
| Oregon | 1324 |
| Indiana | 1164 |
| Minnesota | 1089 |
| South Carolina | 1071 |
| Alabama | 1026 |
| Maryland | 1004 |
| Louisiana | 991 |
| Wisconsin | 990 |
| Nevada | 935 |
| Missouri | 832 |
| Connecticut | 812 |
| Oklahoma | 806 |
| Kentucky | 680 |
| Kansas | 611 |
| Iowa | 538 |
| Arkansas | 434 |
| New Hampshire | 405 |
| Utah | 353 |
| Nebraska | 340 |
| West Virginia | 332 |
| Vermont | 322 |
| Mississippi | 320 |
| Delaware | 311 |
| Maine | 309 |
| Idaho | 295 |
| New Mexico | 291 |
| Hawaii | 276 |
| Alaska | 244 |
| Montana | 219 |
| Rhode Island | 200 |
| Puerto Rico | 125 |
| Wyoming | 64 |
| South Dakota | 55 |
| North Dakota | 54 |
| National | 36 |
| Virgin Islands | 36 |
| Guam | 14 |
| Northern Mariana Islands | 5 |

American Samoa

3

Name: state, dtype: int64

In []: